Is the VLA Backbone Enough to Read Human Goals? Early Goal Disambiguation from Egocentric Video

Vu Thi Thu Trang
Faculty of [TODO: ...]
VNU University of Engineering and Technology
Hanoi, Vietnam
[TODO: email]

Hoang Van Xiem
Faculty of [TODO: ...]
VNU University of Engineering and Technology
Hanoi, Vietnam
[TODO: email]

***Abstract*—Vision–language–action (VLA) policies such as SmolVLA learn the task goal only from a language instruction. On LIBERO-10, SmolVLA succeeds in 19/100 episodes with the correct instruction and in 0/100 when the instruction is empty. A collaborative robot whose partner does not state the goal therefore needs an *intention module* that reads the goal from what the person does. We build such a module for *early goal disambiguation*: from the first $k \in \{2, 4, 6, 8\}$ frames of an egocentric manipulation step, choose the step's purpose among other purposes from the same video. The study proceeds in stages on 2,503 steps (27 videos, 13 scenes) of EgoIntent, with scene-grouped cross-validation and video-level bootstrap intervals. First, a readout trained on frozen features of SmolVLA's own backbone, SmolVLM2, stays near chance (1.13–1.19× chance MRR, every interval includes 1). Second, a pretrained SigLIP encoder used without training raises 5-way accuracy to 27.5–28.9% (chance 20%). Third, a larger SigLIP encoder combined with subtracting the context of earlier steps in the same video reaches 34.1–38.7%, and the same configuration is selected in all 8 folds. Fourth, evaluated once on a separate test split, this configuration reaches 30.2% versus 21.4% for the base pipeline at $k = 8$. Fifth, answering only the most confident half of the queries adds 6.9–7.9 points. On the robot side, we insert the intention as one extra input token and retrain SmolVLA with an always-empty instruction: given only the intention token, it succeeds in 20/100 episodes, on par with 19/100 for the correct instruction, while a wrong intention gives 0/100. Finally, an anticipation study in real seconds shows that the module predicts the upcoming action above chance up to 20 s ahead, that viewing more than 4 s of video does not help, and that most of this ability comes from recognising the current step: it matches a recognise-then-copy baseline, and an AntGPT-style language model reaches 43.8% on the next step only when given the true action history.**

***Index Terms*—intention prediction, egocentric video, vision–language–action models, human–robot collaboration, zero-shot retrieval**

I. Introduction

A robot that helps a person at a table is useful only if it acts on the person's goal. Vision–language–action (VLA) models [1]–[4] map camera images, a language instruction and the robot state to motor commands, and the goal reaches them *only* through the instruction. We measure how strong this dependence is: SmolVLA on LIBERO-10 succeeds in 19/100 episodes with the correct instruction and in 0/100 with an empty one. In collaboration, people rarely state every intention, so the robot must infer the goal from what it sees and pass it to the policy in place of the missing instruction.

Inferring the goal from video is hard because early motion is ambiguous, and because the steps of one activity share the same scene and objects. In a dough-kneading video, "positioning basin", "add salt to flour" and "forming dough" all show the same hands, bowl and flour.

This paper asks how to build an intention module for a VLA policy and answers it in stages. Each stage poses one question, and its result motivates the next (Table I). Our contributions are:

- **The VLA backbone does not read human goals.** A readout trained on frozen SmolVLM2 features reaches only 1.13–1.19× chance MRR on unseen scenes.
- **A training-free intention module.** SigLIP so400m with causal subtraction of earlier-step video context reaches **30.2%** 5-way accuracy on a test split evaluated once, versus 21.4% for base SigLIP and 20% chance. Its confidence margin supports abstention.
- **An intention token can replace the instruction.** SmolVLA fails without goal information (19/100 vs 0/100); retrained with an empty instruction and the goal given as one intention token, it reaches 20/100, and a wrong intention gives 0/100.
- **An honest anticipation analysis.** Measured in seconds against standard anticipation baselines, the module's look-ahead is mostly recognition of the present; the bottleneck is recognising the current step, not knowing the order of steps.

II. Related Work

Egocentric anticipation and intention. Large egocentric corpora [8], [9] enabled long-term action anticipation. Intention has been modelled as a latent variable conditioning future actions [10], inferred with large language models [11], or used to guide reasoning [12]. These works predict future action labels. We isolate a narrower question that matters for

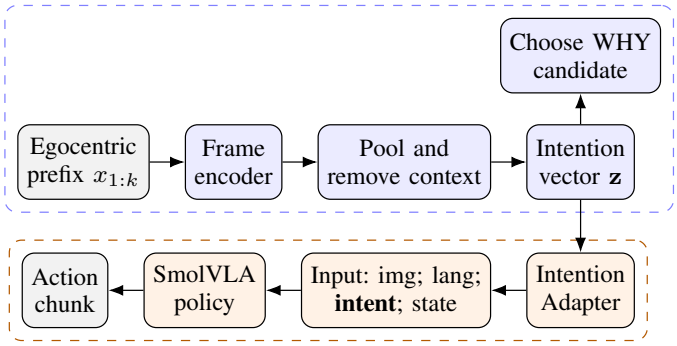


Fig. 1. VLIA. The intention module (top, Stages 1–5, 7–8) maps an egocentric prefix to an intention vector; the adapter turns it into one input token of SmolVLA (bottom, Stage 6).

assistance: discriminating among *competing purposes* inside one activity from a short prefix. EgoIntent [13] provides pre-outcome WHAT/WHY annotations for this.

VLA policies. RT-1/RT-2 [1], [2] and Open X-Embodiment [3] showed that transformer policies on vision–language features generalise across tasks. SmolVLA [4], built on LeRobot [5], couples a 500M SmolVLM2 backbone [6] with a flow-matching [7] action expert. None of these models estimate the intention of a human partner.

Image–text alignment. CLIP [14] and SigLIP [15] learn a joint embedding from web-scale image–caption pairs and transfer zero-shot to new label sets. Our results suggest this alignment is what a small-data intention module lacks.

III. Task, Architecture and Protocol

Task. Each manipulation step i has a video segment ending before its outcome and a free-form purpose sentence y_i (WHY). Eight frames are sampled uniformly over the step, following the sparse segment sampling of TSN [20]; the module sees only the first k. It consists of a video encoder f and a sentence encoder g and scores a candidate sentence c by

$$s(i, c) = \cos\big(f(x_{i,1:k}),\, g(c)\big). \tag{1}$$

The main metric is *5-way accuracy*: choose y_i among itself and 4 other WHY sentences of the *same video* (chance 20%). Distractors are fixed per query, so they are identical across all stages. As a secondary metric we rank y_i against all WHY sentences of the held-out fold (several hundred) and report MRR as a multiple of its exact chance level.

Architecture. In VLIA (Fig. 1) the intention module maps the egocentric prefix to an intention vector $\mathbf{z} \in \mathbb{R}^d$. A linear Intention Adapter projects it to the 960-d hidden size of SmolVLA, and the result is inserted as one token after the instruction and before the state token: $[E_{\text{img}}; E_{\text{lang}}; W\mathbf{z} + b; E_{\text{state}}]$. Because this prefix is cached before the flow-matching steps, the extra token adds negligible cost. Disabling it recovers the original policy exactly.

Data. We develop on the train and validation portions of EgoIntent [13], derived from Ego4D: 2,503 steps, 27 videos, 13 scenes (e.g. kitchen, workshop, farm), 2,032 distinct WHY sentences.

TABLE I
Overview of the stages. 5-way: accuracy of choosing the true purpose among 5 same-video candidates (chance 20%).

Stage	Question	Result	Next step
0	Do structured models help on a small split?	No; 81.5% of errors are same-task	Full data, same-video 5-way
1	Is the VLA backbone enough?	1.13–1.19× chance MRR	Use aligned encoder
2	Does zero-shot alignment help?	5-way 27.5–28.9%	Remove shared context
3	Can training-free choices pass 30%?	5-way 34.1–38.7%, same config 8/8	Freeze, test once
4	Does it hold on unseen videos?	30.2% vs 21.4% ($k = 8$)	Add abstention
5	Can it tell when to wait?	+6.9–7.9 pts at 50% coverage	Connect to policy
6	Can an intention token replace the instruction?	19/100 vs 0/100; token only: 20/100, wrong: 0/100	Predicted intention
7	How far ahead, and how much video?	Plateau from ∼12 s; >4 s of video does not help	Compare baselines
8	Do known improvements help?	No 5-way gain; ≈ recognise-then-copy	Better recognition

TABLE II
Stages 1–2 on scene-held-out folds (2,503 queries). Readout values are means over 3 seeds; brackets are 95% video-bootstrap intervals. Chance: MRR 0.0298, 5-way 0.20.

Encoder	k	MRR	MRR / chance	5-way	5-way / chance
Readout (MiniLM)	2	.0335	1.13 [0.86, 1.43]	.216	1.08 [1.00, 1.16]
	4	.0345	1.16 [0.88, 1.48]	.215	1.07 [1.00, 1.16]
	6	.0353	1.18 [0.90, 1.52]	.212	1.06 [0.98, 1.16]
	8	.0347	1.17 [0.87, 1.51]	.219	1.10 [1.01, 1.20]
Readout (SmolVLM2)	2	.0340	1.14 [0.95, 1.36]	.219	1.09 [1.02, 1.17]
	4	.0338	1.14 [0.94, 1.35]	.222	1.11 [1.03, 1.19]
	6	.0346	1.16 [0.95, 1.39]	.223	1.12 [1.04, 1.20]
	8	.0354	1.19 [0.98, 1.42]	.230	1.15 [1.07, 1.23]
Base SigLIP zero-shot	2	.0671	2.25 [1.87, 2.75]	.277	1.39 [1.23, 1.56]
	4	.0715	2.40 [2.02, 2.88]	.275	1.38 [1.22, 1.55]
	6	.0734	2.47 [2.06, 2.97]	.281	1.40 [1.24, 1.58]
	8	.0765	2.57 [2.13, 3.14]	.289	1.44 [1.27, 1.62]

Evaluation. Scenes are assigned to 8 folds of similar size, and every step is evaluated once, in a scene the model has not seen. Any training or configuration choice uses the training folds only. The EgoIntent test portion (407 steps, 4 videos, after removing one video that overlapped an earlier pilot) is held out and evaluated *once*, for the final configuration. 95% intervals come from 10,000 bootstrap resamples of whole videos; comparisons use paired per-query differences. Pass/fail criteria, including a practical bar of 30% 5-way accuracy, were committed to the repository before each stage.

IV. A Staged Study

A. Stage 0: a pilot on a small split

On 8 videos (556 training and 215 validation steps, one fixed split) with SmolVLM2 frame features, we tried a range of

trained predictors, ranked against all validation goals (random MRR 0.0420). A cosine-only loss collapsed (MRR 0.028, below random); contrastive training with visual input alone gave 0.0540 ± 0.0017. Temporal attention, a Transformer, a GRU and frame deltas all stayed below mean pooling of 8 frames (0.0589 ± 0.0016). Auxiliary WHAT/NEXT heads gave no gain over 3 seeds; feeding predicted WHAT into the WHY path, with same-task hard negatives and a frozen-base reranker, reached 0.0551 ± 0.0009 (an oracle WHAT reached 0.105, a diagnostic only). Ego–exo pretraining on ENIGMA-360 [18] hurt WHY (0.035 vs. 0.048). Most tellingly, 81.5% of top-1 errors picked a goal from *another step of the same task*. With two validation videos and all choices tuned on them, none of these differences was reliable. We therefore moved to the full data with scene-held-out folds, kept mean pooling as the starting point, and added a 5-way test against same-video goals that targets exactly this confusion. The same pilot surfaced a schema bug in our LIBERO runs (6-d instead of 7-d actions); after the fix, the retrained SmolVLA reached 24/100 in its first standard evaluation, and this checkpoint is the one used in Stage 6.

B. Stage 1: a readout on the VLA backbone

Method. Each frame is encoded by frozen SmolVLM2-500M-Video-Instruct (960-d); the first k vectors are mean-pooled and passed through a two-layer MLP (512 hidden units). Targets are a frozen sentence encoder (all-MiniLM-L6-v2 [16]) or SmolVLM2 text features reduced to 256-d by PCA. Training uses multi-positive InfoNCE ($\tau = 0.07$) plus $0.1\times$ a cosine loss, AdamW (lr 10^{-3}), up to 100 epochs chosen on an inner split of the training scenes, three seeds.

Result. The readout is close to chance (Table II): MRR is 1.13–1.19× chance and every interval includes 1; 5-way accuracy is 21.9–23.0% with SmolVLM2 targets. On the original fixed split, validation MRR had looked like 1.7–1.9× chance, but with epochs selected on that same split; on the held-out test portion the readout's best MRR (0.0314) equals chance (0.0318). Letting SmolVLM2 score each candidate by its own likelihood (no training) gives 23.5–24.6% 5-way accuracy, above chance but still weak. *The VLA backbone carries little usable goal signal at this data scale*, which leads to Stage 2.

C. Stage 2: zero-shot image–text alignment

Method. We encode frames with the image tower of `siglip-base-patch16-224` [15] ("base SigLIP"), mean-pool the first k, and encode candidates with the text tower. Nothing is trained.

Result. With no trained parameters, base SigLIP reaches 2.25–2.57× chance MRR; its interval lower bounds (1.87–2.13) lie above the readout's upper bounds at every k. 5-way accuracy is 27.5–28.9%, below the 30% bar. More frames help: from 2 to 8 frames MRR rises by $+0.0094$ [$+0.0053$, $+0.0141$] and 5-way accuracy by $+1.1$ points [$+0.1$, $+2.1$], whereas the readout's gains have intervals containing 0. The typical error recognises the activity but not the step: for "positioning basin", base SigLIP prefers "prepare to knead dough" and "forming cohesive dough". Every frame of the video shows the same table, bowl and hands, and this shared content dominates the embedding. Stage 3 removes it.

TABLE III
5-WAY ACCURACY OF BASE SIGLIP (B) AND THE PROPOSED PIPELINE (P). CV: OUT-OF-FOLD, 2,503 QUERIES. TEST: HELD-OUT SPLIT, 407 QUERIES, 4 VIDEOS, EVALUATED ONCE. Δ: PAIRED P−B IN POINTS, 95% VIDEO-BOOTSTRAP INTERVAL.

	Cross-validation (Stage 3)			Test, run once (Stage 4)		
k	B	P	Δ	B	P	Δ
2	27.7	34.1	+6.4 [3.9, 8.9]	20.4	21.9	+1.5 [−1.0, 3.8]
4	27.5	34.9	+7.4 [4.9, 10.0]	21.1	24.6	+3.4 [1.3, 4.8]
6	28.1	36.2	+8.2 [5.1, 11.2]	22.4	27.8	+5.4 [1.0, 9.7]
8	28.9	38.7	+9.9 [6.2, 13.4]	21.4	30.2	+8.9 [4.4, 14.8]

D. Stage 3: training-free configuration selection

Method. We search 90 training-free configurations: 5 encoders (SigLIP base 224/384, SigLIP2 base, SigLIP so400m, CLIP ViT-L/14) × 3 frame aggregations (mean; max of per-frame cosine; log-sum-exp) × 3 embedding corrections × 2 text templates (raw WHY; "a first-person photo of a person doing: {why}"). The corrections are none, subtracting a global training mean, and *context removal*: subtracting the mean embedding of all frames of *strictly earlier* steps of the same video, which keeps what is new in the current step and stays causal (the first step of a video, 1% of queries, falls back to the global mean). To avoid optimistic selection, each outer fold picks the configuration with the best 5-way accuracy on its *training* folds only and applies it to the held-out fold.

Result. All 8 folds pick the same configuration, *SigLIP so400m, mean pooling, context removal, template*; we call it the *proposed pipeline*. 5-way accuracy rises to 34.1–38.7%, with paired gains of $+6.4$ to $+9.9$ points over base SigLIP, all intervals above 0 (Table III). The MRR gain is positive but its intervals contain 0. Replacing frames by those of a video from another scene ($k = 8$, 313 queries) drops accuracy from 39.6% to 18.8%, so the signal comes from the video. Both parts matter: so400m without context removal gives 31.1–33.0%, base SigLIP with it 29.1–32.1%. The configuration was then frozen into a hashed file before the test portion was read.

E. Stage 4: one evaluation on unseen videos

On the 407 held-out steps, the proposed pipeline is above chance at all four k and above base SigLIP at $k = 4, 6, 8$, reaching 30.2% versus 21.4% at $k = 8$ (Table III). Swapping in another scene's frames drops it from 31.4% to 21.6%. Absolute accuracy is lower than in cross-validation, and base SigLIP drops too, so the test videos appear harder; with only 4 videos the intervals rest on 4 clusters. Gallery MRR of the proposed pipeline is *below* base SigLIP at every k: context removal separates steps within a video, which is the assistance-relevant case, but discards the scene cue that helps ranking across videos.

F. Stage 5: knowing when to wait

Using the top-1 minus top-2 margin as confidence, with thresholds set from the other folds, answering the 50% most confident queries raises accuracy by $+6.9$ to $+7.9$ points (lower bounds ≥ 3.9), and answering 20% reaches 53–61%. On base SigLIP the gain is only 2.6–3.5 points. Three other ideas did not help: stopping observation at the first confident frame did not beat a fixed k at matched frame budgets; mixing in WHY sentences of the nearest training queries lowered MRR; and a hand-region variant was not run because MediaPipe [17] detected hands in only 58.2% of frames (below 30% in four scenes).

G. Stage 6: connecting intention to SmolVLA

Does SmolVLA need the goal? A SmolVLA checkpoint retrained on LIBERO-10 [19] (7-D actions, 25k steps) is run on 10 tasks $\times$ 10 episodes with the correct instruction, an empty one, and another task's instruction; episode e uses the same initial state and seed in every condition. It succeeds in 19/100 episodes with the correct instruction (Wilson [12.5, 27.8]%) and in 0/100 in both other conditions (both [0, 3.7]%); every failed episode runs to the 520-step limit.

Interface. With 256-, 768- and 1152-d intention vectors, the adapter receives gradients, the input grows by exactly one token placed before the state token, the token is present in the KV cache, the flow-matching loss is finite, checkpoints round-trip, and disabling the token reproduces the original SmolVLA. These tests establish correctness, not a behavioural benefit.

Intention instead of instruction. We retrain SmolVLA with the same data and recipe (25k steps) but with the instruction always empty; the task reaches the policy only as an intention vector, the so400m text embedding of the task description (the same encoder and template as the intention module), projected 1152→960 and inserted before the state token. Training was resumed twice from saved checkpoints with optimizer state. Evaluated with an empty instruction on the same 100 (task, episode) units, the correct intention succeeds in **20/100** episodes (Wilson [13.3, 28.9]%), on par with 19/100 for the language policy with the correct instruction (paired difference $+1$ [-7, $+9$] points), while another task's intention gives 0/100 ($+20$ [$+12$, $+28$] points). The tasks it solves are the ones the language policy solves. The intention token therefore carries the goal as well as the instruction did; here it is still the oracle task description, not an intention predicted from video.

H. Stage 7: how far ahead, from how much video?

Method. Robot assistance needs the *upcoming* action. Following the anticipation setting of EPIC-KITCHENS [21], [22], we concatenate the steps of each video, re-sample the 30 fps clips at 4 frames/s, and at anchors every 1 s let the module see the last C seconds (context removal relative to everything before them) and choose the WHAT label of the step that contains $t + H$ among 5 labels of the same video. The same 5,710 anchors (26 videos) and candidates are used for every C and H, so settings are compared with paired differences.

TABLE IV
STAGE 7: 5-WAY ACCURACY (%) OF THE UPCOMING ACTION (WHAT) IN REAL SECONDS; ROWS: SECONDS OF VIDEO SEEN, COLUMNS: SECONDS AHEAD. 5,710 ANCHORS, 26 VIDEOS, CHANCE 20%. LAST ROW: SHARE OF ANCHORS WHOSE TARGET IS STILL THE CURRENT STEP.

Seen	0.5 s	2 s	4 s	8 s	12 s	20 s
1 s	**36.3**	34.3	32.1	29.2	26.2	25.2
2 s	35.9	**34.4**	**32.2**	**30.1**	26.7	25.4
4 s	35.7	34.2	31.9	29.5	26.9	25.6
8 s	34.3	33.2	30.7	28.8	27.0	25.6
16 s	32.7	31.9	31.1	28.8	**27.2**	**26.1**
same step	82%	41%	17%	1%	0%	0%

Result. Accuracy falls from 35.9% at 0.5 s to 26.7% at 12 s and is flat from 12 to 20 s (paired intervals contain 0), still above chance at 20 s (Table IV). Seeing 1, 2 or 4 s is indistinguishable in 30 of 32 paired comparisons, while 8–16 s *hurts* short-horizon prediction by 1–3 points: mean pooling mixes the current action with past ones. The original choice of 8 frames over one step (≈2.6 s) thus lies in the best range. The confidence margin barely changes with the horizon (0.0375 to 0.0351 on the per-step timeline), so the robot cannot tell from it that it is predicting too far ahead.

Baselines. A text-only prior that never sees the video scores 19–20%, so the candidate sets carry no exploitable bias. A *recognise-then-copy* baseline, which recognises the current step with the same module and picks the candidate closest to it, matches the module at every context and horizon (differences within ± 2 points; on the per-step timeline the module is ahead only at 0.6 and 2.6 s, by 1.4–1.7 points). Copying the *true* current label reaches 41.9% at 2.6 s. On the next-step (NEXT) label, an AntGPT-style two-stage model [11], in which Qwen2.5-VL-3B [23] reads the WHAT history as text, reaches 43.8% with the true history but only 27.3–29.2% with the history recognised from video, below the module's 27.8–32.1%; fusing the two scores does not help (31.0% vs 32.1% at $k = 8$). Within-video recognition of the exact current WHAT is only 4–6% top-1, and this, not knowledge of step order, is the bottleneck.

I. Stage 8: do known improvements help?

All on the same folds, with selection on training folds only. Recency-weighted pooling and fusing scores with and without context removal leave 5-way accuracy unchanged (all intervals contain 0) but raise gallery MRR by 10–21%, recovering the cross-video cue lost by context removal. A residual adapter trained with same-video hard negatives [24] is *worse* by 2.0–4.0 points; selection always picked the strongest regularisation. Qwen2.5-VL-3B answering the 5-way question zero-shot reaches 30.8% ($k = 2$) and 39.8% ($k = 8$), on par with the module at $k = 8$, and fusing it with the module adds 2.2 points only at $k = 8$.

V. DISCUSSION AND LIMITATIONS

The main lesson is a design one: with about two thousand labelled steps, an intention module cannot learn the visual–

linguistic mapping from the VLA backbone's features, while a pretrained aligned encoder provides it for free. Removing the context of earlier steps is a cheap, causal addition that helps exactly in the within-activity setting a collaborative robot faces. On the policy side, the instruction slot is where intention should enter: one token carries the goal as well as the instruction. For anticipation, the results point to recognition of the current step as the component to improve, since an oracle history makes even a small language model a strong next-step predictor. The limitations are substantial. The study covers 27 videos in 13 scenes. The test split has only 4 videos, is not scene-disjoint, and shows a clear drop from cross-validation. Context removal lowers cross-video ranking. WHY sentences are free-form and 44.5% of queries have a near-paraphrase in the same video, which caps exact-match scores. The LIBERO experiments use one training seed, one checkpoint and 10 episodes per task, and the intention given to the policy is the oracle task description: a predicted intention has not yet driven the policy, because LIBERO has no matching human video. The anticipation timeline ignores gaps between step clips, and only mean pooling at one frame rate was tested.

VI. Conclusion and Next Step

A readout on the VLA's frozen backbone is near chance at reading human goals, whereas a training-free image–text pipeline with causal context removal reaches 34–39% 5-way accuracy in cross-validation and 30.2% on a test split evaluated once, with a usable confidence signal. SmolVLA retrained with an empty instruction succeeds as often with one intention token (20/100) as with the correct instruction (19/100). The module anticipates upcoming actions above chance up to 20 s, but mostly by recognising the present. The next steps are to drive the policy with intentions predicted from human video, on data that pairs human video with robot tasks, and to improve recognition of the current step. Code and result files are at https://github.com/sgdfgnhmjnfd/vlia-uet.

References

[1] A. Brohan *et al.*, "RT-1: Robotics transformer for real-world control at scale," arXiv:2212.06817, 2022.

[2] B. Zitkovich *et al.*, "RT-2: Vision-language-action models transfer web knowledge to robotic control," in *Proc. CoRL*, 2023, pp. 2165–2183.

[3] Open X-Embodiment Collaboration, "Open X-Embodiment: Robotic learning datasets and RT-X models," in *Proc. ICRA*, 2024, pp. 6892–6903.

[4] M. Shukor *et al.*, "SmolVLA: A vision-language-action model for affordable and efficient robotics," arXiv:2506.01844, 2025.

[5] R. Cadene *et al.*, "LeRobot: An open-source library for end-to-end robot learning," in *Proc. ICLR*, 2026.

[6] A. Marafioti *et al.*, "SmolVLM: Redefining small and efficient multimodal models," arXiv:2504.05299, 2025.

[7] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le, "Flow matching for generative modeling," in *Proc. ICLR*, 2023.

[8] K. Grauman *et al.*, "Ego4D: Around the world in 3,000 hours of egocentric video," in *Proc. CVPR*, 2022, pp. 18995–19012.

[9] K. Grauman *et al.*, "Ego-Exo4D: Understanding skilled human activity from first- and third-person perspectives," in *Proc. CVPR*, 2024, pp. 19383–19400.

[10] E. Valls Mascaro, H. Ahn, and D. Lee, "Intention-conditioned long-term human egocentric action anticipation," in *Proc. WACV*, 2023, pp. 6048–6057.

[11] Q. Zhao *et al.*, "AntGPT: Can large language models help long-term action anticipation from videos?" in *Proc. ICLR*, 2024.

[12] Q. Chu, H. Zhang, M. Liu, Y. Feng, H. Shi, and L. Nie, "Intention-guided cognitive reasoning for egocentric long-term action anticipation," in *Proc. AAAI*, 2026.

[13] "EgoIntent: An open-ended benchmark for intent understanding at pre-outcome micro-steps in egocentric procedural activities," Hugging Face dataset `py2279105943/EgoIntent`, 2026.

[14] A. Radford *et al.*, "Learning transferable visual models from natural language supervision," in *Proc. ICML*, 2021, pp. 8748–8763.

[15] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, "Sigmoid loss for language image pre-training," in *Proc. ICCV*, 2023, pp. 11975–11986.

[16] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proc. EMNLP-IJCNLP*, 2019, pp. 3982–3992.

[17] F. Zhang *et al.*, "MediaPipe Hands: On-device real-time hand tracking," arXiv:2006.10214, 2020.

[18] F. Ragusa, R. Leonardi, M. Mazzamuto *et al.*, "ENIGMA-360: An ego-exo dataset for human behavior understanding in industrial scenarios," arXiv:2603.09741, 2026.

[19] B. Liu, Y. Zhu, C. Gao, Y. Feng, Q. Liu, Y. Zhu, and P. Stone, "LIBERO: Benchmarking knowledge transfer for lifelong robot learning," in *Proc. NeurIPS*, 2023.

[20] L. Wang *et al.*, "Temporal segment networks: Towards good practices for deep action recognition," in *Proc. ECCV*, 2016, pp. 20–36.

[21] D. Damen *et al.*, "Scaling egocentric vision: The EPIC-KITCHENS dataset," in *Proc. ECCV*, 2018, pp. 720–736.

[22] A. Furnari and G. M. Farinella, "Rolling-unrolling LSTMs for action anticipation from first-person video," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 43, no. 11, pp. 4021–4036, 2021.

[23] S. Bai *et al.*, "Qwen2.5-VL technical report," arXiv:2502.13923, 2025.

[24] J. Robinson, C.-Y. Chuang, S. Sra, and S. Jegelka, "Contrastive learning with hard negative samples," in *Proc. ICLR*, 2021.